

Synthesizing Demonic Graph for Worst Case Performance

Mike He, Yihong Zhang
CSE 507 Presentation

Motivations

- Given a program, measure worst case performance is hard
 - Algorithms that work well on random data may fail on specific forms of input data
 - It's hard to construct input data that can expose program's efficiency bottleneck
 - Heuristics make the problem worse
 - Data generated randomly always yield to certain distributions
- But it's good if we can do so
 - Measure reliability of different heuristics
 - Automated data fuzzer for competitive programming events (ICPC, Codeforces, Google Code Jam, etc.)
 - Compare constant factors for various algorithms under worst case
- Let's take SPFA as an example

SPFA for Single Source Shortest Path (SSSP) Problem

```
dist  $\leftarrow$  MAKE-ARRAY(n, [ $\infty$ ]);  
dist[s]  $\leftarrow$  0;  
for i  $\leftarrow$  1 to |V| - 1 do  
    for {u, v, weight}  $\in$  E do  
        if dist[u] + weight < dist[v] then  
            dist[v]  $\leftarrow$  dist[u] + weight;  
        end  
    end  
end
```


Bellman-Ford Algorithm

```
Q  $\leftarrow$  MAKE-QUEUE();  
dist  $\leftarrow$  MAKE-ARRAY(n, [ $\infty$ ]);  
dist[s]  $\leftarrow$  0;  
Push s into Q;  
while Q is not empty do  
    u  $\leftarrow$  Q.POP();  
    for {u, v, weight}  $\in$  E do  
        if dist[u] + weight < dist[v] then  
            dist[v]  $\leftarrow$  dist[u] + weight;  
            Push v into Q if v is not in Q;  
        end  
    end  
end
```

SPFA

How can we hack it?

- Construct a grid graph



Insight: this will create as many suboptimal paths as possible, and each suboptimal paths are expensive to SPFA.

Heuristics!

- Let's add some heuristics!
- Small Label First (SLF)
 - When pushing elements to the queue, check if the label (i.e., distance) on it is less than the one at the **front** of the queue. If so, push it at the front
- Large Label Last (LLL)
 - When pushing elements to the queue, check if the label (i.e., distance) on it is less than the **average value** of labels in the queue. If so, push it at the front
 - * this is actually a modified version of original LLL (at least the one on Wikipedia, which seems to be an counter-optimization)
- Will grid graphs still work?
 - No, they run in near linear time
- What other hacks we can think about?

Sure...

- Construct a Shortest Path Tree (should be as long as possible)
 - Or simply a linked list
- Randomly add edges that connects two vertices with weights slightly greater than the path on the tree
- Shuffle the edges

Insight: this will also create many suboptimal paths, and each suboptimal paths have very similar distance to source, which will make the heuristics feel terrible...

But this SPT-based graph can no longer hack the original algorithm, and it still can't hack SLF.


Synthesizing Demonic Graph with SMT Solver

- Why not use a SMT solver?
- Pros
 - Entire process is automated
 - No longer need sophisticated human insights and empirical experiments
 - No longer need to construct graphs manually---sometimes such graph can be hard to constructed
- Cons
 - Can only synthesize a very small instance
 - We can scale it up!
 - Hard to determine what structural information is needed from the synthesized data
 - There is not an algorithm that is general enough to precisely capture the desired property of arbitrary graphs (More insight, more precision)

Algorithm and Implementation

- Synthesize the gadget
 - Symmetry breaking on graph representation
 - Incremental solving
 - Symbolic-friendly optimization for various SPFA implementations
- Scale the gadget up
 - Specify an entry point and an exit point for the gadget
 - Everytime unroll a vertex into a gadget (self-similar)
 - Different exit point can cause Sager to simulate different structural information
 - We use two versions of start/exit: (start, end) and (start, start), where start and end are the first and last element pushed into the queue.

Scaling: Connect By Exit Nodes



Scale Node 3

Results - 100,000 Nodes & ~300,000 Edges

Graphgen \ Algorithm	SPFA	SPFA+SLF	SPFA+LLL	Dijkstra
Random Graph	≈0.47M	≈0.5M	≈0.6M	≈0.4M
Grid	≈230M	≈0.5M	≈0.5M	≈0.4M
Linked List	≈11M	≈2.2M	≈13M	≈0.7M
Shortest Path Tree	≈2.5M	≈4M	≈ 334M	≈0.4M
Sager	≈ 9,778M	≈ 846M	≈0.8M	≈0.7M

Results - 10,000 Nodes & ~30,000 Edges

Graphgen \ Algorithm	SPFA	SPFA+SLF	SPFA+LLL
Random Graph	≈45K	≈48K	≈58K
Grid	≈1,500K	≈52K	≈49K
Linked List	≈789K	≈223K	≈889K
Shortest Path Tree	≈187K	≈288K	≈ 2M
Sager	≈ 97.8M	≈ 8.6M	≈79k

Limitations

- Sager can only hack short-sighted heuristics like SLF; it's still very hard to synthesize a macro level property that can hack far-sighted heuristics like LLL
- There are some structural properties Sager can't scale
 - E.g., can't scale a grid graph from a square
 - But at least, if you have insight, Sager can help; If you don't, Sager can provide a bottom line!